
EXPERIMENT - 4

AIM:

Develop a game-playing agent using minimax and alpha-beta pruning for Tic-Tac-Toe or a similar turn-based game.

OBJECTIVE:

To develop a game-playing agent for Tic-Tac-Toe using the Minimax algorithm with Alpha-Beta pruning, and to observe how pruning reduces the number of game-tree nodes evaluated while preserving the same optimal decision.

PROCEDURE:

1. Represent the Tic-Tac-Toe board as a 9-cell list and implement a function to detect a winner across all rows, columns and diagonals.
2. Implement the Minimax recursive evaluation function that assigns +10/-10 (adjusted by depth) for an AI win/loss and 0 for a draw at terminal states.
3. Extend Minimax with Alpha-Beta pruning: maintain alpha (best already guaranteed for the maximizer) and beta (best already guaranteed for the minimizer), and prune a branch when $\beta \leq \alpha$.
4. Implement `best_move()` which tries every empty cell for the AI and picks the move with the highest minimax score.
5. Count `nodes_evaluated` on each AI turn to measure the effect of pruning as the game tree shrinks.
6. Simulate a full game: the AI (X) plays optimally via Minimax while the opponent (O) follows a fixed scripted sequence of moves.
7. Print the board after every move along with the number of nodes evaluated, and report the final result (win/loss/draw).

PROGRAM / CODE:

```
# =====
# PRACTICAL 4
# Minimax with Alpha-Beta Pruning - Tic-Tac-Toe
# =====
HUMAN = "O"
AI = "X"
EMPTY = " "
def print_board(board):
    rows = []
    for r in range(3):
        rows.append(" | ".join(board[r * 3:r * 3 + 3]))
    print("\n" + "-" * 9 + "\n").join(rows)
def winner(board):
    lines = [
        (0, 1, 2), (3, 4, 5), (6, 7, 8),
        (0, 3, 6), (1, 4, 7), (2, 5, 8),
        (0, 4, 8), (2, 4, 6)
    ]
    for a, b, c in lines:
        if board[a] != EMPTY and board[a] == board[b] == board[c]:
            return board[a]
    return None
def is_full(board):
    return EMPTY not in board
nodes_evaluated = 0
def minimax(board, depth, alpha, beta, maximizing):
    global nodes_evaluated
    nodes_evaluated += 1
    win = winner(board)
    if win == AI:
        return 10 - depth
```



```

if win == HUMAN:
    return depth - 10
if is_full(board):
    return 0
if maximizing:
    best = float("-inf")
    for i in range(9):
        if board[i] == EMPTY:
            board[i] = AI
            score = minimax(board, depth + 1, alpha, beta, False)
            board[i] = EMPTY
            best = max(best, score)
            alpha = max(alpha, best)
            if beta <= alpha:
                break
    return best
else:
    best = float("inf")
    for i in range(9):
        if board[i] == EMPTY:
            board[i] = HUMAN
            score = minimax(board, depth + 1, alpha, beta, True)
            board[i] = EMPTY
            best = min(best, score)
            beta = min(beta, best)
            if beta <= alpha:
                break
    return best
def best_move(board):
    best_score = float("-inf")
    move = None
    for i in range(9):
        if board[i] == EMPTY:
            board[i] = AI
            score = minimax(board, 0, float("-inf"), float("inf"), False)
            board[i] = EMPTY
            if score > best_score:
                best_score = score
                move = i
    return move
def scripted_human_move(board):
    """A fixed sequence representing a human/opponent playing sensibly."""
    for i in [4, 0, 2, 6, 8, 1, 3, 5, 7]:
        if board[i] == EMPTY:
            return i
if __name__ == "__main__":
    board = [EMPTY] * 9
    turn = AI
    move_no = 0
    print("=" * 55)
    print("  MINIMAX + ALPHA-BETA PRUNING - TIC-TAC-TOE")
    print("=" * 55)
    while True:
        move_no += 1
        print(f"\nMove {move_no} ({turn}'s turn)")
        if turn == AI:
            nodes_evaluated = 0
            pos = best_move(board)
            board[pos] = AI
            print(f"AI plays position {pos} (nodes evaluated: {nodes_evaluated})")
        else:
            pos = scripted_human_move(board)
            board[pos] = HUMAN
            print(f"Opponent plays position {pos}")
        print_board(board)
        w = winner(board)
        if w:
            print(f"\nResult: {w} WINS!")
            break
        if is_full(board):
            print("\nResult: DRAW")
            break
        turn = HUMAN if turn == AI else AI

```



```
print("=" * 55)
```

OUTPUT:

```
=====
MINIMAX + ALPHA-BETA PRUNING - TIC-TAC-TOE
=====
Move 1 (X's turn)
AI plays position 0    (nodes evaluated: 34202)
X |  | 
---
|  | 
---
|  | 
Move 2 (O's turn)
Opponent plays position 4
X |  | 
---
| 0 | 
---
|  | 
Move 3 (X's turn)
AI plays position 1    (nodes evaluated: 1630)
X | X | 
---
| 0 | 
---
|  | 
Move 4 (O's turn)
Opponent plays position 2
X | X | 0
---
| 0 | 
---
|  | 
Move 5 (X's turn)
AI plays position 6    (nodes evaluated: 97)
X | X | 0
---
| 0 | 
---
X |  | 
Move 6 (O's turn)
Opponent plays position 8
X | X | 0
---
| 0 | 
---
X |  | 0
Move 7 (X's turn)
AI plays position 3    (nodes evaluated: 10)
X | X | 0
---
X | 0 | 
---
X |  | 0
Result: X WINS!
=====
```

RESULT:

The Minimax algorithm with Alpha-Beta pruning was successfully implemented for Tic-Tac-Toe. On the AI's first move (empty board, the largest possible search tree), Minimax evaluated 34,202 nodes even with pruning; as the game progressed and the board filled up, the number of nodes evaluated per move fell sharply (1,630 → 97 → 10) because there were fewer empty cells and more branches were pruned early once a guaranteed win/loss bound was found. Playing against a fixed opponent sequence, the AI played optimally at every turn and won the game, confirming that Alpha-Beta pruning returns exactly the same decision as plain Minimax while exploring a fraction of the game tree, which is



essential for making game-tree search practical even in a game as small as Tic-Tac-Toe.